

Vaishnav Gaikwad

Platform & Infrastructure Engineer



Abu Dhabi, UAE | vaishnavgaikwad721@gmail.com | +971 521415768 |
linkedin.com/in/vaishnav-gaikwad | github.com/doflamingo721

Summary

Skilled and experienced Platform/Infrastructure Engineer with a strong focus on Kubernetes, GPU-accelerated AI/ML infrastructure, and the Cloud-Native ecosystem. With five years of hands-on experience, I specialize in deploying and operating large-scale GPU clusters, enabling AI/ML models serving workloads on bare-metal and cloud Kubernetes platforms, and building the observability and security foundations that production AI infrastructure demands. I have deep expertise in NVIDIA GPU stack management - including GPU Operator, Multi-Instance GPU (MIG), DCGM telemetry, and high-performance interconnects - across air-gapped and sovereign cloud environments. I am an individual with great adaptability, problem-solving, collaboration and communication skills with attention to detail, and I always strive to stay ahead of the curve through continuous learning.

Skills

- **Cloud & Container Platforms:** Kubernetes Ecosystem, OpenShift Container Platform, Azure Kubernetes Service, Elastic Kubernetes Service, RKE, RKE2, Rancher, Helm, Docker, Podman, containerd, VMware vSphere (Intermediate), Microsoft Azure, AWS, Google Cloud (Novice), G42 (Novice).
- **AI/ML & GPU Infrastructure:** NVIDIA GPU Operator (CDI mode, precompiled drivers, ClusterPolicy management), Multi-Instance GPU (MIG) configuration and enablement, DCGM Exporter for GPU telemetry, H100 and H200 GPU node management, NVLink readiness and tensor parallelism planning, AI/ML model serving on Kubernetes (vLLM, LLM inference workloads including Gemma, BGE embeddings, TTS/STT models), GPU resource scheduling via device plugins and resource quotas.
- **Kubernetes Networking & Security:** Istio Service Mesh (Intermediate), Kubernetes Network Policies, OPA Gatekeeper, NGINX Ingress Controller, NGINX App Protect, F5 CIS, Kong Ingress, Prisma Cloud, Container Runtime Security, CIS Benchmarks, Compliance Operator.
- **Monitoring & Observability:** Prometheus, Alertmanager, Thanos, Grafana, Loki, DCGM Exporter, Fluentd, ELK Stack, Vector, NGINX Instance Manager.
- **Storage:** ODF/Ceph, MinIO, Portworx CSI, G42 CSI, vSphere CSI.
- **Databases & Messaging:** SQL, MongoDB, Redis, Qdrant (vector database), RabbitMQ, CloudNativePG.
- **Infrastructure as Code & DevOps:** Terraform, Helm, Git, GitHub Actions, GitLab, ArgoCD, Fleet, HashiCorp Vault (VSO, Kubernetes Auth), Red Hat Satellite, Red Hat IdM/LDAP, Red

- Hat Quay, JFrog Artifactory, JFrog Xray, Jira, Confluence.
- **Programming & Scripting:** Bash, Go, Python (Novice), SQL. Previous experience: Java, NodeJS.

Work Overview

Project(s) experience @ Core42 (12/2025 - present)

- **Platform & Infrastructure Engineer:**
 - Deployed and operated OpenShift Container Platform (OCP) clusters in fully air-gapped, bare-metal environments; managed day-2 operations including cluster upgrades, operator lifecycle management via OLM, and compliance hardening using the Compliance Operator.
 - Provisioned and maintained Azure Kubernetes Service (AKS) clusters for cloud workloads, implementing node pool management, cluster autoscaling, Azure Key Vault CSI integration, and managed identity-based secret rotation.
 - Configured NVIDIA GPU infrastructure for H100 and H200 nodes including driver installation, NVIDIA GPU Operator deployment with CDI mode, precompiled drivers, and DCGM Exporter integration for GPU telemetry and observability via Prometheus and Grafana.
 - Deployed and managed AI/ML model serving infrastructure on GPU-enabled Kubernetes clusters; onboarded LLM inference workloads (vLLM, Gemma, BGE embeddings, TTS/STT models) and vector search services (Qdrant), ensuring GPU resource allocation, isolation, and scheduling via device plugins and resource quotas.
 - Deployed and managed Rancher v2 for centralised multi-cluster Kubernetes management in air-gapped environments; configured GitOps workflows using ArgoCD for application lifecycle management across RKE2 clusters.
 - Built and maintained full-stack observability using Prometheus, Grafana, Alertmanager, Loki, and DCGM Exporter; configured ODF/Ceph and MinIO S3-compatible storage backends for stateful workloads, model artifact storage, and message broker infrastructure (RabbitMQ).
 - Managed Red Hat ecosystem tooling including Red Hat Satellite for disconnected content management and patching, and Red Hat IdM for identity and LDAP integration across cluster services.

Project(s) experience @ Huco (08/2023 - 11/2025)

- **SRE & Devops Consultant:**
 - Designed, deployed, and operated secure, air-gapped and public cloud OpenShift and RKE2 clusters, adhering to CIS benchmarks and enterprise hardening best practices.
 - Built and maintained end-to-end CI/CD pipelines for containerized applications using GitOps tools and integrated artifact scanning via JFrog Xray.
 - Configured and managed enterprise-grade ingress and L7 traffic management solutions including F5 CIS, NGINX Ingress Controller, standalone NGINX, and App

- Protect.
- Implemented and managed Istio service mesh, Kubernetes network policies, and Prisma Cloud for container runtime security and threat prevention.
- Deployed and operated observability stacks with Prometheus, Grafana, Thanos, Alertmanager, Fluentd, and ELK, ensuring high availability, scalability, and SLO compliance.
- Provisioned dynamic and static storage with G42, vSphere, Portworx CSI drivers; developed backup and DR strategies for production-grade Kubernetes clusters.
- Established RBAC policies, namespace segregation, and cluster-level governance; assisted development teams with app onboarding and operational support.
- Delivered high availability Rancher setups for central management of RKE2 clusters and led DevOps adoption initiatives for clients embracing cloud-native platforms.

Project(s) experience @ Infracloud Technologies (07/2021 - 08/2023)

Suse Rancher (08/2021 - 08/2023)

- SRE working in the Kubernetes ecosystem. Helping Rancher integrate monthly kubernetes patches and minor versions.
- Maintain and upgrade multiple helm charts supported by the customer.
- Help end users and open-source community members using rancher and RKE with their issues related to provisioning and tools like Istio, Gatekeeper and CIS-benchmark.

Projects

- **Rancher Kubernetes Engine (RKE & RKE2) -**
 - Undertook the responsibilities to add support for the new docker versions(23.0) to RKE. With this Rancher, RKE, RKE2 and K3s users were able to provision kubernetes clusters on machines having docker version 23.0.
 - Introduced multiple docker versions to the rancher's docker installation shell scripts which is used to install docker on different OSes.
 - Integrated monthly kubernetes patches and minor versions with RKE and Rancher. This includes analyzing the changelog of each minor and patch version and implementing these changes in RKE respectively.
 - Added support for newer versions of nginx-ingress-controller and calico/canal CNIs to RKE.
 - Removed multiple deprecated legacy features which were tightly coupled in Rancher's code to optimize the codebase.
 - Resolved multiple longstanding issues related to CNIs and k8s conformance test.
 - Maintaining, upgrading and improving helm charts for Istio, Kiali, Jaeger-tracing, OPA-Gatekeeper and CIS benchmark.
 - Resolved any issues logged by the customers and open source community related to these tools and helm charts.
 - Got hands-on experience with cluster and container security tools like Trivy, kube-bench, sonobuoy while working on the cis-benchmark chart.

Educational Qualification

- **Bachelor of Engineering in Computer Science (2018-2021):**
 - Completed my Bachelor's degree in Computer science from Zeal Institute of Engineering and Research.
- **Diploma in Computer Science (2015-2018):**
 - Completed my Diploma in Computer Science from Government Polytechnic Pune.
- **SSC (2014-2015):**
 - Completed my SSC from Holy Cross Convent High School Akola in 2015.